

ASSIGNMENT 1 & 2

1.Reverse a Given Number

Take the value of the integer and store in a variable, using a while loop, get each digit of the number and store the reversed number in another variable and print the reverse of the number.

```
a=int(input('Enter value: '))
s=0
m=a
while a>0:
    s=int(s*10+(a%10))
    a=int(a/10)
print("Reverse of",m,"is",s)
```

```
➞ Enter value: 523
Reverse of 523 is 325
```

2.Print largest permutation number of a given number

```
def largest_permutation(number):
    digits=[]
    n=number
    while n>0:
        digit=n%10
        digits=[digit]+digits
        n//= 10
    for i in range(len(digits)):
        for j in range(len(digits)-1):
            if digits[j]<digits[j+1]:
                digits[j],digits[j+1]=digits[j+1],digits[j]
    largest_number = 0
    for digit in digits:
        largest_number=largest_number* 10+digit
    return largest_number
number=int(input('Enter value: '))
print("Largest permutation of", number, "is:", largest_permutation(number))
```

```
➞ Enter value: 324
Largest permutation of 324 is: 432
```

3.Find the number of ones in the binary representation of a number

```
a=int(input('Enter value: '))
s=0
c=0
while a>0:
    r=int(a%2)
    if r==1:
        c=c+1
    a=int(a/2)
print(c)
```

```
➞ Enter value: 3
2
```

4. Write a program to print following patterns

a)

```
rows=int(input('Enter range: '))
for i in range(1,rows+1):
    for _ in range(rows-i):
        print(" ", end="")
    for _ in range(i):
        print("**", end="")
    print()
```

```
➞ Enter range: 5
*
**
***
****
*****
```

4. b)

```
rows=int(input('Enter range: '))
for i in range(1,rows+1):
    for _ in range(rows-i):
        print(" ", end="")
    for _ in range(i):
        print("* ", end="")
    print()
```

```
Enter range: 5
      *
     **
    ***
   ****
  *****
 *****
```

4. c)

```
rows=int(input('Enter range: '))
for i in range(1,rows+1):
    for _ in range(i):
        print(chr(64+i), end="")
    print()
```

```
Enter range: 5
A
BB
CCC
DDDD
EEEE
```

5. Check if two numbers are amicable numbers

```
def d_sum(n):
    divisors_sum = 0
    for i in range(1,n):
        if n%i==0:
            divisors_sum+=i
    return divisors_sum
def amicable(num1, num2):
    sum1=d_sum(num1)
    sum2=d_sum(num2)
    return sum1==num2 and sum2==num1
a=int(input("Enter the first number: "))
b=int(input("Enter the second number: "))
if amicable(a, b):
    print("The numbers", a, "and", b, "are amicable.")
else:
    print("The numbers", a, "and", b, "are not amicable.")
```

```
Enter the first number: 220
Enter the second number: 284
The numbers 220 and 284 are amicable.
```

6. Find the cumulative sum of a list where the i-th element is the sum of the first i+1 elements from the original list.

```
a=input("Enter the list elements separated by spaces: ")
a=[int(x) for x in a.split()]
ans=[]
total=0
for i in a:
    total+=i
    ans.append(total)
print("Original List:", a)
print("Cumulative Sum List:", ans)
```

```
Enter the list elements separated by spaces: 1 2 3 4 5
Original List: [1, 2, 3, 4, 5]
Cumulative Sum List: [1, 3, 6, 10, 15]
```

[+ Code](#)
[+ Text](#)

7. Given a list of sorted numbers and a variable K, where K is also a number, write a Python program using binary search to find the number in the list which is closest to the given number K

```
n = input("Enter the sorted list of numbers separated by spaces: ")
sorted_list = [int(x) for x in n.split()]
k = int(input("Enter the number K: "))
low = 0
high = len(sorted_list) - 1
ans = None
```

```

ans = None
while low <= high:
    mid = (low + high) // 2
    if sorted_list[mid] == k:
        ans = sorted_list[mid]
        break
    elif sorted_list[mid] < k:
        low = mid + 1
    else:
        high = mid - 1
if ans is None or abs(sorted_list[mid] - k) < abs(ans - k):
    ans = sorted_list[mid]
print("Number in the list closest to", k, "is:", ans)

```

```

➦ Enter the sorted list of numbers separated by spaces: 2 5 7 8 12
Enter the number K: 9
Number in the list closest to 9 is: 8

```

8. Given a list of tuples, write a Python program to remove all the duplicated tuples from the given list using the concept of set.

```

t_list=[(1, 2), (3, 4), (1, 2), (5, 6), (3, 4)]
t_set=set(t_list)
t_list1=list(t_set)
print("Original list of tuples:",t_list)
print("List of unique tuples:", t_list1)

```

```

➦ Original list of tuples: [(1, 2), (3, 4), (1, 2), (5, 6), (3, 4)]
List of unique tuples: [(1, 2), (3, 4), (5, 6)]

```

9. Given an unsorted list of some elements (may or may not be integers), Find the frequency of each distinct element in the list using a dictionary.

```

def frequency(a):
    f_dict = {}
    for i in a:
        if i in f_dict:
            f_dict[i]+=1
        else:
            f_dict[i]=1
    return f_dict
unsorted_list = [1, 2, 1, 2, 1, 'a','b','a','a']
f=frequency(unsorted_list)
print("Frequency of each distinct element is:")
for i,freq in f.items():
    print(i, ":", freq)

```

```

➦ Frequency of each distinct element is:
1 : 3
2 : 2
a : 3
b : 1

```

10. Given two words, check whether they are anagrams using dictionary.

```

def anagram(a,b):
    a=a.lower()
    b=b.lower()
    if len(a)!=len(b):
        return False
    f1={}
    f2={}
    for char in a:
        if char in f1:
            f1[char]+=1
        else:
            f1[char]=1
    for char in b:
        if char in f2:
            f2[char]+=1
        else:
            f2[char]=1
    return f1==f2
a=input("Enter the first word: ")
b=input("Enter the second word: ")
if anagram(a,b):
    print(f"{a} and {b} are anagrams.")
else:
    print(f"{a} and {b} are not anagrams.")

```

↩ Enter the first word: moon
Enter the second word: mono
moon and mono are anagrams.

11. Find common elements in three sorted lists using sets.

```
common = lambda a, b, c: set(a) & set(b) & set(c)
list1 = [1, 2, 3, 4, 5, 8]
list2 = [2, 4, 6, 8, 10]
list3 = [3, 4, 7, 8, 9]
ans = common(list1, list2, list3)
print("Common elements in the three lists:", ans)
```

↩ Common elements in the three lists: {8, 4}

12. Find Symmetric Pairs in dictionary using loop.

```
def symmetric_pair(d):
    ans=[]
    for key,value in d.items():
        if value in d and d[value]==key:
            ans.append((key, value))
    return ans
d={'a':'b', 'b':'a', 'c':'d', 'd':'e', 'e':'d'}
ans=symmetric_pair(d)
print("Symmetric pairs in the dictionary:",ans)
```

↩ Symmetric pairs in the dictionary: [('a', 'b'), ('b', 'a'), ('d', 'e'), ('e', 'd')]

KNN CLASSIFICATION IMPLEMENTATION

CODE

```
import pandas as pd
dataset = pd.read_csv("/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/iris.csv")
print(dataset.head())
X = dataset.iloc[:, 1:5]
y = dataset.iloc[:, 5]
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
#from sklearn.preprocessing import StandardScaler
#scaler = StandardScaler()
#scaler.fit(X)
#X_train = scaler.transform(X_train)
#X_test = scaler.transform(X_test)
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors = 5)
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
import pickle
output_model_file = 'Knnmodel.pkl'
with open(output_model_file, 'wb') as f:
    pickle.dump(classifier, f)
```

OUTPUT

```
↗ Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm Species
0 1 5.1 3.5 1.4 0.2 Iris-setosa
1 2 4.9 3.0 1.4 0.2 Iris-setosa
2 3 4.7 3.2 1.3 0.2 Iris-setosa
3 4 4.6 3.1 1.5 0.2 Iris-setosa
4 5 5.0 3.6 1.4 0.2 Iris-setosa
[[ 6 0 0]
 [ 0 10 0]
 [ 0 2 12]]
precision recall f1-score support
Iris-setosa 1.00 1.00 1.00 6
Iris-versicolor 0.83 1.00 0.91 10
Iris-virginica 1.00 0.86 0.92 14
accuracy 0.93 30
macro avg 0.94 0.95 0.94 30
weighted avg 0.94 0.93 0.93 30
```


```
from google.colab import drive
drive.mount('/content/drive')
```

```
↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

CODE

```
import numpy as np
error = []
for i in range(1, 10):
    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(X_train, y_train)
    pred_i = knn.predict(X_test)
    error.append(np.mean(pred_i != y_test))
import matplotlib.pyplot as plt
plt.plot(range(1, 10), error,
         color='red', linestyle='dashed',
         marker='o', markerfacecolor='blue',
         markersize=10)
plt.title('Error Rate K Value')
plt.xlabel('K Value')
plt.ylabel('Mean Error')
plt.show()
```

OUTPUT



KNN CLASSIFICATION IMPLEMENTATION

CODE

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

data = pd.read_csv("/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/Titanic.csv")
x = data.drop('Survived', axis = 1)
y = data['Survived']
x.drop(['Name', 'Ticket', 'Cabin'], axis = 1, inplace = True)
print(data.shape)
print(data.isna().sum())
# Missing value Imputation
# data = data.dropna(axis = 0, how = 'any')
# Checking for missing value
# print(data.isna().sum())
# print(data.shape)
# numeric value imputation with mean
x['Age'] = x['Age'].fillna(x['Age'].mean())
x['Embarked'] = x['Embarked'].fillna(x['Embarked'].mode()[0])
x = pd.get_dummies(x, columns = ['Sex', 'Embarked'], prefix = ['Sex', 'Embarked'], drop_first = True)
print(x.head())
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state = 0)
from sklearn.preprocessing import StandardScaler
std_x = StandardScaler()
x_train = std_x.fit_transform(x_train)
x_test = std_x.transform(x_test)
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors = 5)
classifier.fit(x_train, y_train)
y_pred = classifier.predict(x_test)
from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

OUTPUT

```
(891, 12)
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch           0
Ticket           0
Fare             0
Cabin           687
Embarked         2
dtype: int64
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             0
SibSp            0
Parch           0
Ticket           0
Fare             0
Cabin            0
Embarked         0
dtype: int64
(183, 12)
  PassengerId  Pclass  Age  SibSp  Parch  Fare  Sex_male  Embarked_Q \
0            1      3  22.0      1      0   7.2500      1          0
1            2      1  38.0      1      0  71.2833      0          0
2            3      3  26.0      0      0   7.9250      0          0
3            4      1  35.0      1      0  53.1000      0          0
4            5      3  35.0      0      0   8.0500      1          0

  Embarked_S
0           1
1           0
2           1
3           1
4           1
[[99 11]
 [19 50]]
      precision    recall  f1-score   support

         0       0.84      0.90      0.87        110
         1       0.82      0.72      0.77         69

   accuracy
macro avg       0.83      0.81      0.82        179
weighted avg       0.83      0.83      0.83        179
```

NAIVE BAYESIAN CLASSIFICATION IMPLEMENTATION

CODE

```
# Assigning features and label variables
weather=['Sunny','Sunny','Overcast','Rainy','Rainy','Rainy','Overcast','Sunny','Sunny',
'Rainy','Sunny','Overcast','Overcast','Rainy']
temp=['Hot','Hot','Hot','Mild','Cool','Cool','Cool','Mild','Cool','Mild','Mild','Mild','Hot','Mild']

play=['No','No','Yes','Yes','Yes','No','Yes','No','Yes','Yes','Yes','Yes','Yes','No']

# Import LabelEncoder
from sklearn import preprocessing
#creating labelEncoder
le = preprocessing.LabelEncoder()
# Converting string labels into numbers.
weather_encoded=le.fit_transform(weather)
temp_encoded=le.fit_transform(temp)
target=le.fit_transform(play)

print(weather_encoded)
print(temp_encoded)
print(target)

import numpy as np
zipped=zip(weather_encoded,temp_encoded)
features = np.array(list(zipped)).tolist()
print(features)

#Import Gaussian Naive Bayes model
from sklearn.naive_bayes import CategoricalNB

#Create a Gaussian Classifier
model = CategoricalNB()

# Train the model using the training sets
model.fit(features,target)

#Predict Output
predicted= model.predict([[0, 2]]) # 0:Overcast, 2:Mild
print("Predicted Value:", predicted)
```

OUTPUT

```
↗ [2 2 0 1 1 1 0 2 2 1 2 0 0 1]
  [1 1 1 2 0 0 0 2 0 2 2 2 1 2]
  [0 0 1 1 1 0 1 0 1 1 1 1 1 0]
  [[2, 1], [2, 1], [0, 1], [1, 2], [1, 0], [1, 0], [0, 0], [2, 2], [2, 0], [1, 2], [2, 2], [0, 2], [0, 1], [1, 2]]
  Predicted Value: [1]
```

DECISION TREE CLASSIFICATION IMPLEMENTATION

CODE

```
# Assigning features and label variables
weather=['Sunny','Sunny','Overcast','Rainy','Rainy','Rainy','Overcast','Sunny','Sunny',
'Rainy','Sunny','Overcast','Overcast','Rainy']
temp=['Hot','Hot','Hot','Mild','Cool','Cool','Cool','Mild','Cool','Mild','Mild','Mild','Hot','Mild']

play=['No','No','Yes','Yes','Yes','No','Yes','No','Yes','Yes','Yes','Yes','Yes','No']

# Import LabelEncoder
from sklearn import preprocessing
#creating labelEncoder
le = preprocessing.LabelEncoder()
# Converting string labels into numbers.
weather_encoded=le.fit_transform(weather)
temp_encoded=le.fit_transform(temp)
target=le.fit_transform(play)

print(weather_encoded)
print(temp_encoded)
print(target)

import numpy as np
zipped=zip(weather_encoded,temp_encoded)
features = np.array(list(zipped)).tolist()
print(features)

from sklearn import tree
#Create a Gaussian Classifier
model = tree.DecisionTreeClassifier(criterion='entropy')

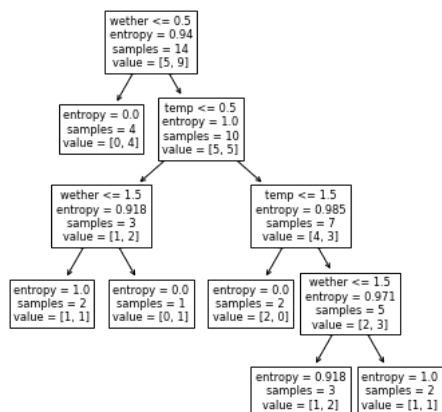
# Train the model using the training sets
model.fit(features,target)

#Predict Output
predicted= model.predict([[0, 2]]) # 0:Overcast, 2:Mild
print("Predicted Value:", predicted)

from matplotlib import pyplot as plt
fig, ax = plt.subplots(figsize=(6, 6)) #figsize value changes the size of plot
tree.plot_tree(model,ax=ax,feature_names=['wether','temp'])
plt.show()
```

OUTPUT

```
[2 2 0 1 1 1 0 2 2 1 2 0 0 1]
[1 1 1 2 0 0 0 2 0 2 2 2 1 2]
[0 0 1 1 1 0 1 0 1 1 1 1 1 0]
[[2, 1], [2, 1], [0, 1], [1, 2], [1, 0], [1, 0], [0, 0], [2, 2], [2, 0], [1, 2], [2,
Predicted Value: [1]
```



LINEAR REGRESSION

```
from google.colab import drive
drive.mount('/content/drive')
```

↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

The aim is to build a model which predicts sales based on the money spent on different platforms such as TV, radio, and newspaper for marketing.

CODE

```
#Importing the libraries import pandas
as pd
import numpy as np
import matplotlib.pyplot as plt import seaborn
as sns
```

```
#Reading the dataset
dataset = pd.read_csv("advertising.csv")
```

```
dataset.head()
```

OUTPUT

↗

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

□ Data Pre-Processing

```
dataset.shape
```

↗ (200, 4)

1. Checking for missing values

```
dataset.isna().sum()
```

OUTPUT

↗

TV	0
Radio	0
Newspaper	0
Sales	0

dtype: int64

Conclusion: The dataset does not have missing values

2. Checking for duplicate rows

```
dataset.duplicated().any()
```

↗ False

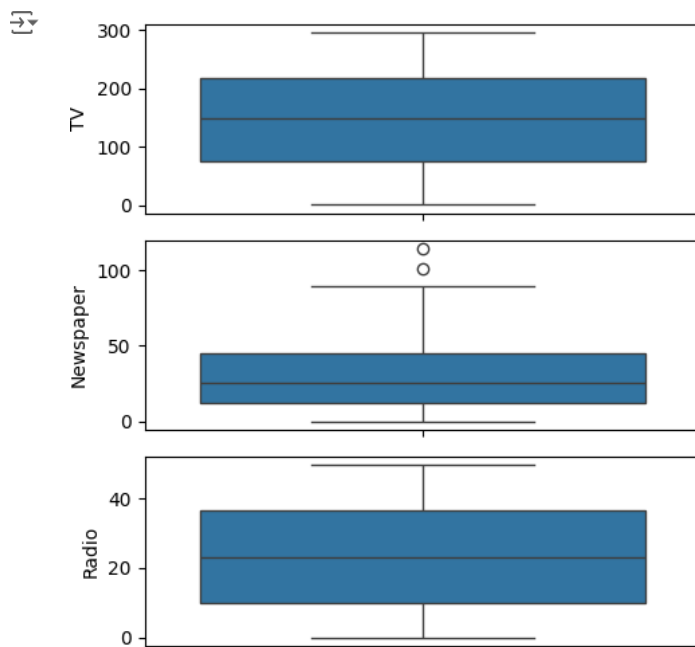
Conclusion: There are no duplicate rows present in the dataset

3. Checking for outliers

CODE

```
fig, axs = plt.subplots(3, figsize = (5,5))
plt1 = sns.boxplot(dataset['TV'], ax = axs[0])
plt2 = sns.boxplot(dataset['Newspaper'], ax = axs[1]) plt3 =
sns.boxplot(dataset['Radio'], ax = axs[2])
plt.tight_layout()
```

OUTPUT



Conclusion: There are not that extreme values present in the dataset

□ Exploratory Data Analysis

1. Distribution of the target variable

CODE

```
sns.distplot(dataset['Sales']);
```

OUTPUT


<ipython-input-11-e26ae89dfd77>:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(dataset['Sales']);
```



Conclusion: It is normally distributed

2. How Sales are related with other variables

```
sns.pairplot(dataset, x_vars=['TV', 'Radio', 'Newspaper'], y_vars='Sales', height=4, aspect=1, kind='scatter') plt.show()
```



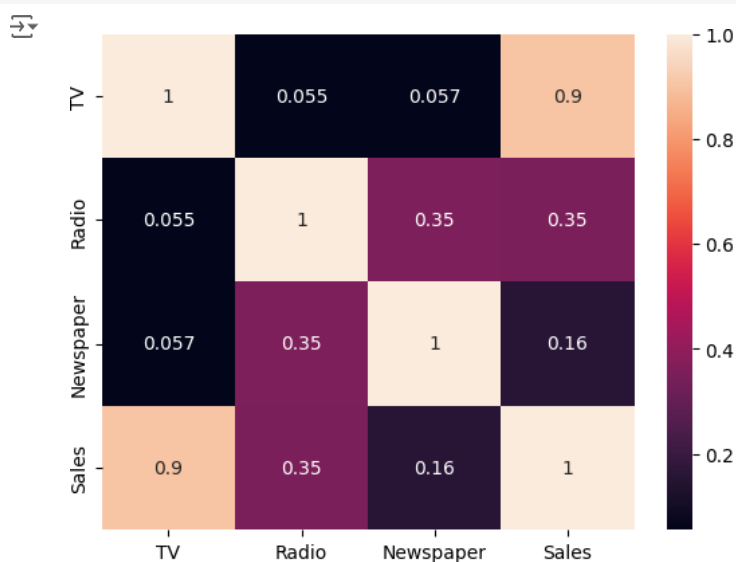
Conclusion: TV is strongly, positively, linearly correlated with the target variable. Bu the Newspaper feature seems to be uncorrelated

3. Heatmap

CODE

```
sns.heatmap(dataset.corr(), annot = True) plt.show()
```

OUTPUT



Conclusion: TV seems to be most correlated with Sales as 0.9 is very close to 1

1. Simple Linear Regression

Simple linear regression has only one x and one y variable. It is an approach for predicting a quantitative response using a single feature.

It establishes the relationship between two variables using a straight line. Linear regression attempts to draw a line that comes closest to the data by finding the slope and intercept that define the line and minimize regression errors.

Formula: $Y = \beta_0 + \beta_1 X + e$

Y = Dependent variable / Target variable β_0 = Intercept

of the regression line

β_1 = Slope of the regression line which tells whether the line is increasing or decreasing

X = Independent variable / Predictor variable e = Error

Equation: Sales = $\beta_0 + \beta_1 X + TV$

CODE

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
#Setting the value for X and Y x =
dataset[['TV']]
y = dataset['Sales']
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.3, random_state = 100)
```

```
slr= LinearRegression()
slr.fit(x_train.values, y_train)
```



LinearRegression

LinearRegression()

```
#Printing the model coefficients
print('Intercept: ', slr.intercept_)
print('Coefficient:', slr.coef_)
```



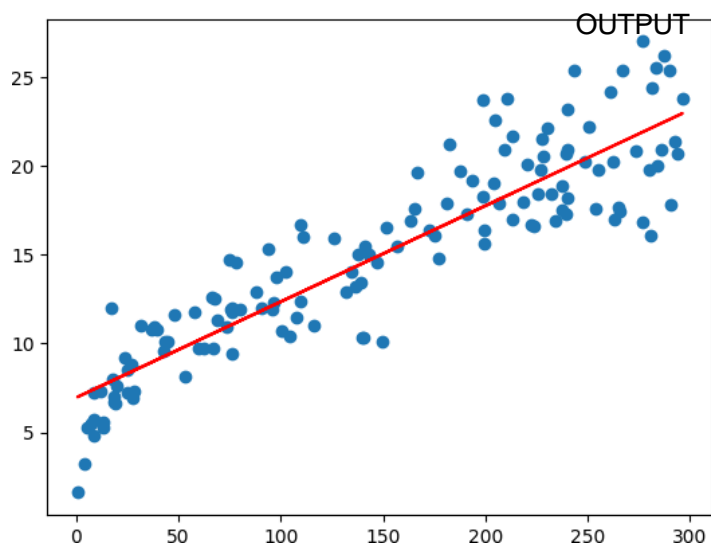
Intercept: 6.948683200001357
Coefficient: [0.05454575]

```
print('Regression Equation: Sales = 6.948 + 0.054 * TV')
```



Regression Equation: Sales = 6.948 + 0.054 * TV

```
#Line of best fit
plt.scatter(x_train, y_train)
plt.plot(x_train, 6.948 + 0.054*x_train, 'r') plt.show()
```



```
#Prediction of Test and Training set result y_pred_slr=
slr.predict(x_test.values)
x_pred_slr= slr.predict(x_train.values)
```

CODE

```
print("Prediction for test set: {}".format(y_pred_slr))
```

OUTPUT



Prediction for test set: [7.37414007 19.94148154 14.32326899 18.82329361 20.13239168 18.2287449
14.54145201 17.72692398 18.75238413 18.77420243 13.34144544 19.46693349
10.01415451 17.1923756 11.70507285 12.08689312 15.11418241 16.23237035
15.8669138 13.1068987 18.65965635 14.00690363 17.60692332 16.60328147
17.03419291 18.96511257 18.93783969 11.05597839 17.03419291 13.66326538
10.6796127 10.71234015 13.5487193 17.22510305 9.67597085 13.52144643
12.25053038 16.13418799 19.07965865 17.48692266 18.69783838 16.53237199
15.92145955 18.86693021 13.5050827 11.84143724 7.87050642 20.51966653
10.79961336 9.03233096 17.99419817 16.29237067 11.04506924 14.09963141
18.44147334 9.3759692 7.88687015 8.34505447 17.72692398 11.62325422]


CODE

```
#Actual value and the predicted value
slr_diff = pd.DataFrame({'Actual value': y_test, 'Predicted value': y_pred_slr}) slr_diff
```


 Show hidden output

```
plt.plot(y_test.values, color = 'red', label = 'Real data')
plt.plot(y_pred_slr, color = 'blue', label = 'Predicted data') plt.title('Prediction')
plt.legend() plt.show()
```

OUTPUT



```
#Predict for any value slr.predict([[56]])
```

 array([10.00324536])

Conclusion: The model predicted the Sales of 10.003 in that market

```
# print the R-squared value for the model from
sklearn.metrics import accuracy_score
print('R squared value of the model: {:.2f}'.format(slr.score(x,y)*100))
```


 R squared value of the model: 81.10
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but LinearRegression was fitted with warnings.warn(

Conclusion: 81.10% of the data fit the regression model

CODE

```
# 0 means the model is perfect. Therefore the value should be as close to 0 as possible meanAbErr =
metrics.mean_absolute_error(y_test, y_pred_slr)
meanSqErr = metrics.mean_squared_error(y_test, y_pred_slr)
rootMeanSqErr = np.sqrt(metrics.mean_squared_error(y_test, y_pred_slr))

print('Mean Absolute Error:', meanAbErr) print('Mean
Square Error:', meanSqErr)
print('Root Mean Square Error:', rootMeanSqErr)

 Mean Absolute Error: 1.6480589869746525 Mean
Square Error: 4.077556371826948
Root Mean Square Error: 2.019296008966231
```

OUTPUT

2. Multiple Linear Regression

Multiple linear regression has one y and two or more x variables. It is an extension of Simple Linear regression as it takes more than one predictor variable to predict the response variable.

Multiple Linear Regression is one of the important regression algorithms which models the linear relationship between a single dependent continuous variable and more than one independent variable.

Assumptions for Multiple Linear Regression: 1. A linear relationship should exist between the Target and predictor variables. 2. The regression

residuals must be normally distributed. 3. MLR assumes little or no multicollinearity (correlation between the independent variable) in data.

Formula: $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \dots + \beta_n X_n + e$

Y = Dependent variable / Target variable β_0 = Intercept

of the regression line

$\beta_1, \beta_2, \dots, \beta_n$ = Slope of the regression line which tells whether the line is increasing or decreasing X_1, X_2, \dots, X_n = Independent variables /

Predictor variables

e = Error

Equation: Sales = $\beta_0 + (\beta_1 * TV) + (\beta_2 * Radio) + (\beta_3 * Newspaper)$

CODE

```
#Setting the value for X and Y
x = dataset[['TV', 'Radio', 'Newspaper']] y =
dataset['Sales']
```

```
x_train, x_test, y_train, y_test= train_test_split(x, y, test_size= 0.3, random_state=100)
```

```
mlr= LinearRegression()
mlr.fit(x_train.values, y_train)
```



▼ LinearRegression

LinearRegression()

```
#Printing the model coefficients
print(mlr.intercept_)
# pair the feature names with the coefficients list(zip(x, mlr.coef_))
```



```
4.334595861728431
[('TV', 0.053829108667250075),
 ('Radio', 0.11001224388558054),
 ('Newspaper', 0.0062899501461303325)]
```

```
#Predicting the Test and Train set result y_pred_mlr=
mlr.predict(x_test.values)
x_pred_mlr= mlr.predict(x_train.values)
```

```
print("Prediction for test set: {}".format(y_pred_mlr))
```

OUTPUT



```
Prediction for test set: [ 9.35221067 20.96344625 16.48851064 20.10971005 21.67148354 16.16054424
 13.5618056  15.39338129 20.81980757 21.00537077 12.29451311 20.70848608
  8.17367308 16.82471534 10.48954832  9.99530649 16.34698901 14.5758119
 17.23065133 12.56890735 18.55715915 12.12402775 20.43312609 17.78017811
 16.73623408 21.60387629 20.13532087 10.82559967 19.12782848 14.84537816
 13.13597397  9.07757918 12.07834143 16.62824427  8.41792841 14.0456697
  9.92050209 14.26101605 16.76262961 17.17185467 18.88797595 15.50165469
 15.78688377 16.86266686 13.03405813 10.47673934 10.6141644  20.85264977
 10.1517568   6.88471443 17.88702583 18.16013938 12.55907083 16.28189561
 18.98024679 11.33714913  5.91026916 10.06159509 17.62383031 13.19628335]
```

CODE


```
#Actual value and the predicted value
mlr_diff = pd.DataFrame({'Actual value': y_test, 'Predicted value': y_pred_mlr}) mlr_diff
```



Show hidden output

OUTPUT

```
plt.plot(y_test.values, color = 'red', label = 'Real data')
plt.plot(y_pred_mlr, color = 'blue', label = 'Predicted data') plt.title('Prediction')
plt.legend() plt.show()
```



```
#Predict for any value  
mlr.predict([[56, 55, 67]])
```

OUTPUT



```
array([13.82112602])
```

Conclusion: The model predicted the Sales of 13.82 in that market

```
# print the R-squared value for the model  
print('R squared value of the model: {:.2f}'.format(mlr.score(x,y)*100))
```



```
R squared value of the model: 90.11  
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but LinearRegression was fitted with warnings.warn(
```

Conclusion: 90.21% of the data fit the multiple regression model

CODE

```
# 0 means the model is perfect. Therefore the value should be as close to 0 as possible  
meanAbsErr = metrics.mean_absolute_error(y_test, y_pred_mlr)  
meanSqErr = metrics.mean_squared_error(y_test, y_pred_mlr)  
rootMeanSqErr = np.sqrt(metrics.mean_squared_error(y_test, y_pred_mlr))  
  
print('Mean Absolute Error:', meanAbsErr) print('Mean  
Square Error:', meanSqErr)  
print('Root Mean Square Error:', rootMeanSqErr)
```

OUTPUT



```
Mean Absolute Error: 1.227818356658941  
Mean Square Error: 2.6360765623280655  
Root Mean Square Error: 1.623599877533891
```

CODE

OUTPUT

CODE

```

import numpy
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import silhouette_score
from sklearn.cluster import DBSCAN

def dbscan(D, eps, MinPts):
    labels = [0]*len(D)
    C = 0
    for P in range(0, len(D)):
        if not (labels[P] == 0):
            continue
        NeighborPts = region_query(D, P, eps)
        else:
            C += 1
            grow_cluster(D, labels, P, NeighborPts, C, eps, MinPts)
    return labels

def grow_cluster(D, labels, P, NeighborPts, C, eps, MinPts):
    labels[P] = C
    i = 0
    while i < len(NeighborPts):
        Pn = NeighborPts[i]
        if labels[Pn] == -1:
            labels[Pn] = C
        elif labels[Pn] == 0:
            labels[Pn] = C
            PnNeighborPts = region_query(D, Pn, eps)
            if len(PnNeighborPts) >= MinPts:
                NeighborPts = NeighborPts + PnNeighborPts
        i += 1

def region_query(D, P, eps):
    neighbors = []
    for Pn in range(0, len(D)):
        if numpy.linalg.norm(D[P] - D[Pn]) < eps:
            neighbors.append(Pn)
    return neighbors

data = pd.read_csv("/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/Mall_Customers.csv")
#f1 = data['Age'].values
f2 = data['Annual Income (k$)'].values
f3 = data['Spending Score (1-100)'].values
X = numpy.array(list(zip(f2, f3)))
cluster_labels = numpy.array(dbscan(X, 3, 4))
#db = DBSCAN(eps = 3, min_samples = 4).fit(X)
#cluster_labels = db.labels_
print(cluster_labels)


n_clusters_ = len(set(cluster_labels)) - (1 if -1 in cluster_labels else 0)
n_noise_ = list(cluster_labels).count(-1)

print("Estimated number of clusters: ", n_clusters_)
print("Estimated number of noise points: ", n_noise_)

plt.figure(figsize=(7,5))
plt.scatter(X[cluster_labels == 0, 0], X[cluster_labels == 0, 1], s = 50, c = 'pink')
plt.scatter(X[cluster_labels == 1, 0], X[cluster_labels == 1, 1], s = 50, c = 'yellow')
plt.scatter(X[cluster_labels == 2, 0], X[cluster_labels == 2, 1], s = 50, c = 'cyan')
plt.scatter(X[cluster_labels == 3, 0], X[cluster_labels == 3, 1], s = 50, c = 'magenta')
plt.scatter(X[cluster_labels == 4, 0], X[cluster_labels == 4, 1], s = 50, c = 'orange')
plt.scatter(X[cluster_labels == 5, 0], X[cluster_labels == 5, 1], s = 50, c = 'blue')
plt.scatter(X[cluster_labels == 6, 0], X[cluster_labels == 6, 1], s = 50, c = 'red')
plt.scatter(X[cluster_labels == 7, 0], X[cluster_labels == 7, 1], s = 50, c = 'black')
plt.scatter(X[cluster_labels == 8, 0], X[cluster_labels == 8, 1], s = 50, c = 'green')
plt.xlabel('Annual Income in (1k)')
plt.ylabel('Spending Score from 1-100')
plt.title('Clusters of data')
plt.show()
if(n_clusters_ > 1):
    sil = silhouette_score(X, cluster_labels,
                           metric='euclidean',
                           sample_size = len(data))
    print("Quality of Clustering: ", sil)

```


Estimated number of noise points: 135




K-MEANS CLUSTERING

CODE


```
from sklearn.cluster import KMeans
from sklearn import metrics
from scipy.spatial.distance import cdist
import numpy as np
import matplotlib.pyplot as plt
x1 = np.array([3, 1, 1, 2, 1, 6, 6, 6, 5, 6, 7, 8, 9, 8, 9, 9, 8])
x2 = np.array([5, 4, 5, 6, 5, 8, 6, 7, 6, 7, 1, 2, 1, 2, 3, 2, 3])
X = np.array(list(zip(x1, x2))).reshape(len(x1), 2)
plt.scatter(X[:,0],X[:,1], label='True Position')
plt.title('Dataset')
plt.scatter(x1, x2)
plt.show()
```


OUTPUT




```
import matplotlib.pyplot as plt
distortions = []
mapping1 = {}
for k in range(1, 10):
    kmeanModel = KMeans(n_clusters=k)
    kmeanModel.fit(X)
    sse = sum(np.min(cdist(X, kmeanModel.cluster_centers_, 'euclidean'), axis=1)) / X.shape[0]
    print("Value of K = ", k, ", SSE = ", sse)
    distortions.append(sum(np.min(cdist(X, kmeanModel.cluster_centers_, 'euclidean'), axis=1)) / X.shape[0])
    mapping1[k] = sum(np.min(cdist(X, kmeanModel.cluster_centers_, 'euclidean'), axis=1)) / X.shape[0]
    plt.scatter(X[:,0],X[:,1], c=kmeanModel.labels_, cmap='rainbow')
    plt.scatter(kmeanModel.cluster_centers_[:,0],kmeanModel.cluster_centers_[:,1], color='black')
    plt.show()
for key, val in mapping1.items():
    print(f'{key} : {val}')
plt.plot(range(1, 10), distortions, 'bx-')
plt.xlabel('Values of K')
plt.ylabel('Distortion')
plt.title('The Elbow Method using Distortion')
plt.show()
```

OUTPUT


 /usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: warnings.warn(
Value of K = 1 , SSE = 3.4577032384495707



/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: warnings.warn(
Value of K = 2 , SSE = 1.7687413573405673




/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: warnings.warn(
Value of K = 3 , SSE = 0.8819889697423957




/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: warnings.warn(
Value of K = 4 , SSE = 0.7587138847606585







```
/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning:
warnings.warn(
Value of K = 5 , SSE = 0.6760729098960964
```




```
/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning:
warnings.warn(
Value of K = 6 , SSE = 0.580097449143775
```




```
/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning:
warnings.warn(
Value of K = 7 , SSE = 0.517480107950963
```






/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: warnings.warn(
Value of K = 8 , SSE = 0.42618267462691206





/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: warnings.warn(
Value of K = 9 , SSE = 0.35294117647058826



1 : 3.4577032384495707
2 : 1.7687413573405673
3 : 0.8819889697423957
4 : 0.7587138847606585
5 : 0.6760729098960964
6 : 0.580097449143775
7 : 0.517480107950963
8 : 0.42618267462691206
9 : 0.35294117647058826

The Elbow Method using Distortion







CODE

```
import matplotlib.pyplot as plt
import numpy as np
X = np.array([[5,3], [10,15], [15,12], [24,10], [30,45], [85,70], [71,80], [60,78], [55,52], [80,91],])
plt.scatter(X[:,0],X[:,1], label='True Position')
plt.show()
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
kmeans = KMeans(n_clusters=2)
kmeans.fit(X)
print(kmeans.cluster_centers_)
print(kmeans.labels_)
cluster_labels = kmeans.predict(X)
C = kmeans.cluster_centers_
sil = silhouette_score(X, cluster_labels, metric='euclidean',sample_size = len(X))
print(C)
print("Quality of Clustering: ", sil)
plt.scatter(X[:,0],X[:,1], c=kmeans.labels_, cmap='rainbow')
plt.scatter(kmeans.cluster_centers_[0],kmeans.cluster_centers_[1], color='black')
plt.show()
```

OUTPUT



```
/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning:
  warnings.warn(
[[70.2 74.2]
 [16.8 17. ]]
[1 1 1 1 0 0 0 0 0]
[[70.2 74.2]
 [16.8 17. ]]
Quality of Clustering: 0.6586004781412067
```



CODE

```
import pandas as pd
import numpy as np
```

```

import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
data = pd.read_csv("/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/spinem.csv")
X = data[['pelvic_incidence', 'pelvic_radius', 'thoracic_slope']]
kmeans = KMeans(n_clusters = 3, random_state = 123)
model = kmeans.fit(X)
cluster_labels = kmeans.predict(X)
X['Cluster'] = cluster_labels
print(X)
C = kmeans.cluster_centers_
sil = silhouette_score(X, cluster_labels, metric='euclidean', sample_size = len(data))
print(C)
print("Quality of Clustering: ", sil)
fig = plt.figure()
plt.scatter(X['pelvic_incidence'], X['pelvic_radius'], c=cluster_labels,
            s=50, cmap='viridis');
plt.scatter(C[:, 0], C[:, 1], marker='*', s=1000)
fig = plt.figure()
ax = Axes3D(fig)
ax.scatter(X['pelvic_incidence'], X['pelvic_radius'],
            X['thoracic_slope'],
            c=cluster_labels,
            cmap='viridis');
ax.scatter(C[:, 0], C[:, 1], C[:, 2],
            marker='*',
            c='#050505')

```

OUTPUT

```

/usr/local/lib/python3.9/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning:
warnings.warn(

```

	pelvic_incidence	pelvic_radius	thoracic_slope	Cluster
0	63.027817	98.672917	14.5386	2
1	39.056951	114.405425	17.5323	0
2	68.832021	105.985135	17.4861	2
3	69.297008	101.868495	12.7074	2
4	49.712859	108.168725	15.9546	2
..
305	47.903565	117.449062	14.7484	0
306	53.936748	114.365845	18.1972	0
307	61.446597	125.670725	13.5565	0
308	45.252792	118.545842	16.0928	0
309	33.841641	123.945244	17.6963	0

```

[310 rows x 4 columns]
[[ 46.42903837 124.47018491 13.26250567]
 [ 80.49567418 120.00557969 12.78133516]
 [ 62.59438009 103.64870812 13.03697051]]

```

Quality of Clustering: 0.3539586349354204

```

<ipython-input-1-cecf937cbdde>:12: SettingWithCopyWarning:

```


A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/us>

```

X['Cluster'] = cluster_labels
<mpl_toolkits.mplot3d.art3d.Path3DCollection at 0x7febd0c27790>

```



<Figure size 640x480 with 0 Axes>

The following code block implements k-means algorithm from the scratch

CODE


```

from sklearn.metrics import pairwise_distances_argmin
def find_clusters_kmeans(X, n_clusters, rseed=2):
    centers = [X[0], X[1], X[2]]
    print(centers)
    while True:
        labels = pairwise_distances_argmin(X, centers)
        new_centers = np.array([X[labels == i].mean(0)
                                for i in range(n_clusters)])
        if np.all(centers == new_centers):
            break
        centers = new_centers
    return centers, labels
X = np.array(X)
centers, cluster_labels = find_clusters_kmeans(X, 3)
plt.scatter(X[:, 0], X[:, 1], c=cluster_labels,
            s=50, cmap='viridis');
sil = silhouette_score(X, cluster_labels, metric='euclidean', sample_size = len(data))
print(centers)
print("Quality of Clustering: ", sil)


```

OUTPUT

```

[array([63.0278175 , 98.67291675, 14.5386      , 2.
        ], array([ 39.05695098, 11
        ]), array([ 62.59438009 103.64870812 13.03697051 2.
        ], array([ 46.42903837 124.47018491 13.26250567 0.
        ], array([ 80.49567418 120.00557969 12.78133516 1.
        ])]
Quality of Clustering: 0.3539586349354204

```



APRIORI ALGORITHM (MARKET BASKET ANALYSIS)

CODE

```
1 dataset = [['Milk', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],
2           ['Dill', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],
3           ['Milk', 'Apple', 'Kidney Beans', 'Eggs'],
4           ['Milk', 'Unicorn', 'Corn', 'Kidney Beans', 'Yogurt'],
5           ['Corn', 'Onion', 'Onion', 'Kidney Beans', 'Ice cream', 'Eggs']]
```

```
1 import pandas as pd
2 from mlxtend.preprocessing import TransactionEncoder
3 te = TransactionEncoder()
4 te_ary = te.fit(dataset).transform(dataset)
5 df = pd.DataFrame(te_ary, columns=te.columns_)
```

OUTPUT



	Apple	Corn	Dill	Eggs	Ice cream	Kidney Beans	Milk	Nutmeg	Onion	Unicorn	Yogurt
0	False	False	False	True	False	True	True	True	True	False	True
1	False	False	True	True	False	True	False	True	True	False	True
2	True	False	False	True	False	True	True	False	False	False	False
3	False	True	False	False	False	True	True	False	False	True	True
4	False	True	False	True	True	True	False	False	True	False	False

CODE

```
1 from mlxtend.frequent_patterns import apriori
2 frequent_itemsets = apriori(df, min_support=0.6, use_colnames=True)
3 frequent_itemsets['length'] = frequent_itemsets['itemsets'].apply(lambda x: len(x))
4 frequent_itemsets
```

OUTPUT



	support	itemsets	length
0	0.8	(Eggs)	1
1	1.0	(Kidney Beans)	1
2	0.6	(Milk)	1
3	0.6	(Onion)	1
4	0.6	(Yogurt)	1
5	0.8	(Eggs, Kidney Beans)	2
6	0.6	(Onion, Eggs)	2
7	0.6	(Milk, Kidney Beans)	2
8	0.6	(Onion, Kidney Beans)	2
9	0.6	(Yogurt, Kidney Beans)	2
10	0.6	(Onion, Eggs, Kidney Beans)	3

CODE

```
1 frequent_itemsets[(frequent_itemsets['length'] > 2) &
2                   (frequent_itemsets['support'] >= 0.6)]
```



	support	itemsets	length
10	0.6	(Onion, Eggs, Kidney Beans)	3

CODE

```
1 from mlxtend.frequent_patterns import association_rules
2
3 association_rules(frequent_itemsets, metric="confidence", min_threshold=0.7)
```

OUTPUT



	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction
0	(Eggs)	(Kidney Beans)	0.8	1.0	0.8	1.00	1.00	0.00	inf
1	(Kidney Beans)	(Eggs)	1.0	0.8	0.8	0.80	1.00	0.00	1.0
2	(Onion)	(Eggs)	0.6	0.8	0.6	1.00	1.25	0.12	inf
3	(Eggs)	(Onion)	0.8	0.6	0.6	0.75	1.25	0.12	1.6
4	(Milk)	(Kidney Beans)	0.6	1.0	0.6	1.00	1.00	0.00	inf
5	(Onion)	(Kidney Beans)	0.6	1.0	0.6	1.00	1.00	0.00	inf
6	(Yogurt)	(Kidney Beans)	0.6	1.0	0.6	1.00	1.00	0.00	inf
7	(Onion, Eggs)	(Kidney Beans)	0.6	1.0	0.6	1.00	1.00	0.00	inf
8	(Onion, Kidney Beans)	(Eggs)	0.6	0.8	0.6	1.00	1.25	0.12	inf
9	(Eggs, Kidney Beans)	(Onion)	0.8	0.6	0.6	0.75	1.25	0.12	1.6
10	(Onion)	(Eggs, Kidney Beans)	0.6	0.8	0.6	1.00	1.25	0.12	inf

IMAGE CLASSIFICATION USING KNN, SVC, ANN

CODE

```
import cv2
import os
```

```
def extract_color_histogram(image):
    hsv = cv2.cvtColor(image, cv2.COLOR_BGR2HSV)
    bins = [10]
    h, s, v = hsv[:, :, 0], hsv[:, :, 1], hsv[:, :, 2]
    hist = cv2.calcHist([h], [0], None, bins, [0, 180])
    cv2.normalize(hist, hist)
    return hist.flatten()
```

```
train_data_path = '/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/ImageData/FruitData/Training'
test_data_path = '/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/ImageData/FruitData/Test'
data_dir_list = list(os.listdir(train_data_path))
print(data_dir_list)
features = []
classLabels = []
for dataset in data_dir_list:
    img_list = os.listdir(train_data_path+'/'+dataset)
    print ('Loaded the images of dataset-'+'{ }\n'.format(dataset))
    for img in img_list:
        image = cv2.imread(train_data_path + '/' + dataset + '/' + img )
        label = dataset
        hist = extract_color_histogram(image)
        features.append(hist)
        classLabels.append(label)
```

OUTPUT

```
['Apple', 'Banana']
Loaded the images of dataset-Apple

Loaded the images of dataset-Banana
```

CODE

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
trainFeat, testFeat, trainLabels, testLabels = train_test_split(features, classLabels, test_size=0.20)
```

```
from sklearn.neighbors import KNeighborsClassifier
print("\n")
print("[INFO] evaluating k-NN...")
k = 9
model = KNeighborsClassifier(n_neighbors = k)
model.fit(trainFeat, trainLabels)
acc = model.score(testFeat, testLabels)
print("[INFO] k-NN classifier: k = ", k)
print("[INFO] accuracy: {:.2f}%".format(acc * 100))
predLabels = model.predict(testFeat)
print(confusion_matrix(testLabels, predLabels))
print(classification_report(testLabels, predLabels))
test_img_list = os.listdir(test_data_path)
for img in test_img_list:
    print(test_data_path + '/' + img)
    image = cv2.imread(test_data_path + '/' + img )
    hist = extract_color_histogram(image)
    prediction = model.predict([hist])
    print("Predicted Class Label = ",prediction)
```

OUTPUT

```
[INFO] evaluating k-NN...
[INFO] k-NN classifier: k = 9
[INFO] accuracy: 100.00%
[[100  0]
 [ 0 97]]
      precision    recall  f1-score   support

   Apple         1.00      1.00      1.00        100
   Banana         1.00      1.00      1.00         97

 accuracy          1.00          1.00          1.00        197
 macro avg          1.00          1.00          1.00        197
 weighted avg          1.00          1.00          1.00        197
```

```
/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/ImageData/FruitData/Test/banana_test.jpg
Predicted Class Label = ['Banana']
/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/ImageData/FruitData/Test/apple_test.jpg
Predicted Class Label = ['Apple']
```

CODE

```
from sklearn.svm import SVC
print("\n")
print("[INFO] evaluating SVC...")
model = SVC(max_iter=1000, class_weight='balanced')
model.fit(trainFeat, trainLabels)
acc = model.score(testFeat, testLabels)
print("[INFO] SVC classifier")
print("[INFO] accuracy: {:.2f}%".format(acc * 100))
predLabels = model.predict(testFeat)
print(confusion_matrix(testLabels, predLabels))
print(classification_report(testLabels, predLabels))
test_img_list = os.listdir(test_data_path)
for img in test_img_list:
    print(test_data_path + '/' + img)
    image = cv2.imread(test_data_path + '/' + img)
    hist = extract_color_histogram(image)
    prediction = model.predict([hist])
    print("Predicted Class Label = ", prediction)
```

OUTPUT



```
[INFO] evaluating SVC...
[INFO] SVC classifier
[INFO] accuracy: 100.00%
[[100  0]
 [ 0  97]]
      precision    recall  f1-score   support

   Apple         1.00      1.00      1.00        100
   Banana         1.00      1.00      1.00         97

 accuracy         1.00      1.00      1.00        197
 macro avg         1.00      1.00      1.00        197
 weighted avg         1.00      1.00      1.00        197

/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/ImageData/FruitData/Test/banana_test.jpg
Predicted Class Label = ['Banana']
/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/ImageData/FruitData/Test/apple_test.jpg
Predicted Class Label = ['Apple']
```

CODE

```
from sklearn.neural_network import MLPClassifier
#Neural Network
print("\n")
print("[INFO] evaluating ANN...")
model = MLPClassifier(hidden_layer_sizes=(10, 10), max_iter=1000, solver='sgd', learning_rate_init=.1)
model.fit(trainFeat, trainLabels)
acc = model.score(testFeat, testLabels)
print("[INFO] Neural Network accuracy: {:.2f}%".format(acc * 100))
predLabels = model.predict(testFeat)
print(confusion_matrix(testLabels, predLabels))
print(classification_report(testLabels, predLabels))
test_img_list = os.listdir(test_data_path)
for img in test_img_list:
    print(test_data_path + '/' + img)
    image = cv2.imread(test_data_path + '/' + img)
    hist = extract_color_histogram(image)
    prediction = model.predict([hist])
    print("Predicted Class Label = ", prediction)
```

OUTPUT



```
[INFO] evaluating ANN...
[INFO] Neural Network accuracy: 100.00%
[[100  0]
 [ 0  97]]
      precision    recall  f1-score   support

   Apple         1.00      1.00      1.00        100
   Banana         1.00      1.00      1.00         97

 accuracy         1.00      1.00      1.00        197
 macro avg         1.00      1.00      1.00        197
 weighted avg         1.00      1.00      1.00        197

/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/ImageData/FruitData/Test/banana_test.jpg
Predicted Class Label = ['Banana']
/content/drive/MyDrive/JISNIT/Courses/ML/LectureNotes/data/ImageData/FruitData/Test/apple_test.jpg
Predicted Class Label = ['Apple']
```